Mixed-Initiative Coordination for Disaster Response in the Real-World

Paper XXX

ABSTRACT

The problem of allocating emergency responders to rescue tasks is a key application area for agent-based coordination algorithms. However, to date, none of the proposed approaches take into account the uncertainty predominant in disaster scenarios and, crucially, none have been deployed in the real-world in order to understand how humans perform when instructed by an agent. Hence, in this paper, we propose a novel algorithm, using Multi-agent Markov Decision Processes to coordinate emergency responders. More importantly, we deploy this algorithm in a mixed-reality game to help an agent guide human players to complete rescue tasks. In our field trials, our algorithm is shown to improve human performance and our results allow us to elucidate some of the key challenges faced when deploying of mixed-initiative team formation algorithms.

Categories and Subject Descriptors

H.4 [**Information Systems Applications**]: Multi-Agent Systems

General Terms

Design, Human Factors, Algorithms

Keywords

Human-Agent Interaction, Coordination, Decision under Uncertainty, Adjustable Autonomy

1. INTRODUCTION

The coordination of teams of field responders in search and rescue missions is regarded as one of the grand challenges for multi-agent systems research [**?**]. In such settings, responders with different capabilities (e.g., fire extinguishing, digging, or life support) have to form teams in order to perform rescue tasks (e.g., extinguishing a fire or digging civilians out of rubble or both) to minimise costs (e.g., time or money) and maximise the number of lives and buildings saved. Thus, responders have to plan their paths to the tasks (as these may be distributed in space) and form specific teams to complete some tasks. These teams, in turn,

may need to disband and reform other teams to complete other tasks requiring different capabilities, taking into account the status (e.g., health or building fire) of the tasks and the environment (e.g., if a fire or radioactive cloud is spreading). Furthermore, uncertainty in the environment (e.g., wind direction or speed) or in the responders' abilities to complete tasks (e.g., some may be tired or get hurt) means that plans may need to change depending on the state of the players and the environment.

To address these challenges, in recent years, a number of algorithms and mechanisms have been developed to create teams and allocate tasks. For example, [**?**, **?**, **?**] and [**?**, **?**], developed centralised and decentralised optimisation algorithms respectively to allocate rescue tasks efficiently to teams of field responders with different capabilities. However, none of these approaches considered the uncertainty in the environment or in the field responders' abilities. Crucially, to date, while all of these algorithms have been shown to perform well in simulations (assuming agents as computational entities), none of them have been *deployed* to guide *real* human responders (amateur or expert) in real-time rescue missions. Thus, it is still unclear whether these algorithms will cope with real-world uncertainties (e.g., communication breakdowns or change in wind direction), be acceptable to humans (i.e., agent-computed plans are not confusing and take into account human capabilities), and do help humans perform better than on their own.

Against this background, in this paper we develop a novel algorithm for team coordination under uncertainty and evaluate it within a real-world mixed-reality game that embodies the simulation of team coordination in disaster response settings. In more detail, we consider scenario involving rescue tasks distributed in a disaster space over which a radioactive cloud is spreading. Tasks need to be completed by the responders before the area is completely covered by the cloud (as responders will die from radiation exposure) which is spreading according to varying wind speed and direction. Our algorithm captures the uncertainty in the scenario (i.e., in terms of environment and player states) and is able to compute a policy to allocate responders to tasks to minimise the time to complete all tasks without them being exposed to significant radiation. The algorithm is then used by an agent to guide human responders based on their perceived states. This agent is then implemented in our deployed platform, AtomicOrchid, that structures the interaction between human responders, a human coordinator, and the agent in a mixed-reality location-based game. By so doing, we are able to study, both quantitatively and qualitatively, the perfor-

mance of a mixed-initiative team (i.e., a human team under human and agent guidance) and the interactions between the different actors in the system. Thus, this paper advances the state of the art in the following ways:

1. We develop a novel approximate algorithm for team formation under uncertainty using a Multi-agent Markov Decision Process (MMDP) paradigm, and show how it accounts for real-world uncertainties.

2. We present AtomicOrchid, a novel platform to evaluate team formation under uncertainty using the concept of mixed-reality games. AtomicOrchid allows an agent, using our team formation algorithm, to coordinate, in real-time, human players using mobile phone-based messaging, to complete rescue tasks efficiently.

3. We provide the first real-world evaluation of a team formation agent in a disaster response setting in field trials and present both quantitative and qualitative results. Our results allow us to elucidate some of the challenges for the formation of human-agent collectives, that is, mixed-initiative teams where control can be passed between agents and humans in flexible ways.

When taken together, our results show, for the first time, how agent-based coordination algorithms for disaster response can be validated in the real-world. Moreover, these results allow us to derive a methodology and guidelines to evaluate human-agent interaction in real-world settings.

The rest of this paper is structured as follows. Section **??** formalises the disaster response problem as an MMDP. Section **??** then describes the algorithm to solve the path planning and task allocation problems presented by the MMDP while Section **??** describes the AtomicOrchid platform. Section **??** presents our pilot study and the evaluation of the system in a number of field trials. Finally, Section **??** concludes.

2. THE DISASTER SCENARIO

We consider a disaster scenario involving a satellite, powered radioactive fuel, that has crashed in a sub-urban area (see Section **??** to see how this helps implement a mixed-reality game). While debris is strewn around a large area, damaging buildings and causing accidents and injuring civilians, radioactive particles discharged in the air, from the debris, are gradually spreading over the area, threatening to contaminate food reserves and people. Hence, emergency services, voluntary organisations, and the military are deployed to help evacuate the casualties and resources before these are engulfed by a radioactive cloud. In what follows, we model this scenario formally and then describe the optimisation problem faced by the actors (i.e., including emergency services, volunteers, medics, and soldiers) in trying to save as many lives and resources as possible. We then propose an algorithm to solve this optimisation problem (in Section **??**). In Section **??**, we demonstrate how this algorithm can be used by a software agent (in our mixed-reality game) in a mixed-initiative process to coordinate field responders.

2.1 Formal Model

Let G denote a grid overlaid on top of the disaster space, and the satellite and actors are located at various coordinates $(x, y) \in G$ in this grid. The radioactive cloud induces a radioactivity level $l \in [0, 100]$ at every point it covers in the grid (100 corresponds to maximum radiation). While the exact radiation levels can be measured by responders on the ground (at every grid coordinate) using their geiger counter, it is assumed that some information is available from existing sensors in the area. However, this information is uncertain due to the poor positioning of the sensors and the variations in wind speed and direction (and we show how this uncertainty is captured in the next section). A number of safe zones $G' \subseteq G$ are defined where the responders can drop off assets and casualties. Let the set of n field responders be denoted as $p_1, \cdots p_i \cdots, p_n \in I$ and the set of ms rescue tasks as $t_1, \cdots, t_j, \cdots, t_m \in T$. As responders enact tasks, they may become tired, get injured, or receive radiation doses that may, at worst, be life threatening. Hence, we assign each responder a health level $h_i \in [0, 100]$ that decreases based on their radiation dose (**Wenchao: what's the relationship between radiation dose and health decrease?**) and assume that their decision to perform the task allocated to them is liable to some uncertainty (e.g., they may not want to do a task because they are tired or don't believe it is the right one to do). Moreover, each responder has a specific role $r \in Roles$ (e.g., fire brigade, soldier, or medic) and this will determine the capabilities he or she has and therefore the tasks he or she can perform. We denote as $Roles(p_i)$ the role of responder p_i. In turn, to complete a given task t_j, a set of responders $I' \subseteq I$ with specific roles $R_{t_j} \subseteq R$ is required. Thus, a task can only be completed by a team of responders I' if $\{Roles(p_i)|p_i \in I'\} = R_{t_j}$. In our scenario, we assume
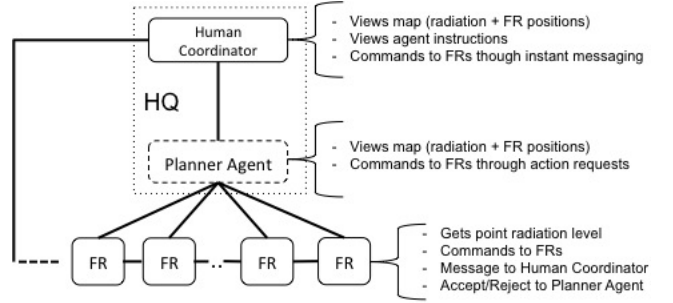

Figure 1: The interactions between different actors in the disaster scenario. Lines represent communication links. Planner agent and coordinator sit in the headquarters (HQ). Field responders (FRs) can communicate with all actors directly.

that the field responders are coordinated by the headquarters headed by a human coordinator H but assisted by a planner agent PA that can receive input from the human coordinator or the field responders. While the human coordinator H communicates its instructions directly to the responders (e.g., using an instant messaging client or walkie talkie), the planning agent PA can compute an allocation of tasks for the responders to complete. This is communicated to them in terms of simple "Do task X at position Y". The responders may not want to do some tasks (for reasons outlined above) and may therefore simply accept or reject the received instruction. These interactions are depicted in Figure 2.1.

It is important to note that our model implements different types of control: (i) agent-based: when the agent instructs the responders (ii) human-based: when responders work with the coordinator or between themselves. Our model also captures different modes of control: (i) centralised: when responders respond to the planning agent or human coordinator (ii) decentralised: if responders coordinate between themselves. Crucially, this scenario allows for flexible levels of human and agent autonomy. For example, field responders may simply implement the plan given to them by the planner agent but can also feedback their constraints to the planner agent (as we demonstrate later) by rejecting some instructions and requesting new instructions.

Given this model, we next formulate the optimisation problem faced by the responders and solved by the planning agent (later in Section **??**). To this end, we propose a Multi-Agent Markov Decision Process (MMDP) [**?**] that captures the uncertainties of the radioactive cloud and the responders' behaviours. Specifically, we model the spreading of the radiative cloud as a random process over the disaster space and allow the actions requested from the responders to fail (because they refuse to go to a task) or incur delays (because they are too slow) during the rescue process. This stands in contrast to previous work [**?, ?**] that require the process of task executions to be deterministic and explicitly model the task deadlines as deterministic constraints (which are stochastic in our domain). Thus in the MMDP model, we represent task executions as stochastic processes of state transitions. Thus, the uncertainties of the radioactive cloud and the responders' behaviours can be easily captured with transition probabilities. Additionally, modelling the problem as a MMDP allows us to use many sophisticated algorithms such as RTDP and MCTS that have already been well developed in the literature [**?, ?, ?**].

2.2 The Optimisation Problem

Here we formalise the optimisation problem that needs to be solved to coordinate the responders optimally. Hence, we define this problem as a Multi-agent Markov Decision Process (MMDP) formally represented by tuple $\mathcal{M} = \langle I, S, \{A_i\}, P, R \rangle$, where I is the set of actors as defined in the previous section, S is the state space, A_i is a set of responder p_i's actions, P is the transition function, and R is the reward function. We elaborate on each of these below.

More specifically, $S = S_r \times S_{p_1} \times \cdots \times S_{p_n} \times S_{t_1} \times \cdots \times S_{t_m}$ where $S_r = \{l_{(x,y)} | (x,y) \in G\}$ is the state variable of the radioactive cloud to specify the radioactive level $l_{(x,y)} \in [0,100]$ at every point $(x,y) \in G$. $S_{p_i} = \langle h_i, (x_i, y_i), t_j \rangle$ is the state variable for each responder p_i to specify his or her health level $h_i \in [0,100]$, the coordinate (x_i, y_j), and the task t_j carried by the responder. $S_{t_j} = \langle st_j, (x_j, y_j) \rangle$ is then the state variable for task t_j to specify its status st_j (picked up, dropped off, or idle) and coordinate (x_j, y_j).

The three types of actions (in set A_i) a responder can take are: (i) *stay* in the current location (x_i, y_i), (ii) *move* to the 8 neighbouring locations, or (iii) *complete* a task located in (x_i, y_i). A joint action $\vec{a} = \langle a_1, \cdots, a_n \rangle$ is a set of actions where $a_i \in A_i$, one for each responder.

The transition function P is defined in more detail as: $P = P_r \times P_{p_1} \times P_{p_n} \times P_{t_1} \times P_{t_n}$ where:

- $P_r(s'_r | s_r)$ is the probability for the radioactive cloud to spread from state s_r to s'_r. It captures the uncertainty of the next radioactive levels of the environment due to the noisy sensor reading and the variation in wind speed and direction.
- $P_{p_i}(s'_{p_i} | s, a_i)$ is the probability for responder p_i to transit to a new state s'_{p_i} when executing action a_i. For example, when a responder is asked to go to a new location, he or she may not be there because he or she becomes tired, gets injured, or receives radiation doses that are life threatening.
- $P_{t_j}(s'_{t_j} | s, \vec{a})$ is the probability for task t_j. A task t_j can only be completed by a team of responders with required roles locating in the same coordinate as t_j.

Now, if a task t_j is completed (i.e., in s_{t_j}, the status st_j is marked as "dropped off" and the coordinate (x_j, y_j) is within a drop off zone), the team will be rewarded using function R. There will be a penalty for the team if a responder p_i gets injured or receives a high dose of radiation (i.e., in s_{p_i}, the health level h_i is 0). Moreover, we attribute a cost to each of the responders' since it will requires them to exert some effort (e.g., running or carrying objects).

Give the above definitions, a policy for the MMDP is a mapping from states to joint actions, $\pi : S \to \vec{A}$ so that the responders know which actions to take given the current state of the problem. The quality of a policy π is usually measured by its expected value V^π, which can be computed recursively by the Bellman equation:

$$V^\pi(s) = R(s, \pi(s)) + \gamma \sum_{s' \in S} P(s' | s, \pi(s)) V^\pi(s') \quad (1)$$

where $\pi(s)$ is a joint action given s and $\gamma \in (0,1]$ is the discounted factor. The goal of solving the MMDP is to find an optimal policy π^* that maximises the expected value with the initial state s^0, $\pi^* = \arg\max_\pi V^\pi(s^0)$.

At each decision step, we assume the planning agent can fully observe the state of the environment s by collecting sensor readings of the radioactive cloud and GPS locations of the responders. Given a policy π of the MMDP, a joint action $\vec{a} = \pi(s)$ can be selected and broadcast to the responders (as mentioned earlier). By so doing, each responder can be instructed by the agent and know how to act in the field. In the next section we discuss the computational challenges of finding an optimal policy and propose a scalable approximation algorithm for this purpose.

3. TEAM COORDINATION ALGORITHM

The MMDP model proposed in the previous section results in a very large search space even for small-sized problems. For example, with 8 responders and 17 tasks in a 50×55 grid, the number of possible states is more than 2×10^{400}. Therefore, it is practically impossible to compute the optimal solution. In such cases, it is therefore better to consider approximate solution approaches that result in high quality allocations. The point of departure for our approximate solution comes from the observations that, when making a decision, the responders first need to *cooperatively* select a task to form a team with others (i.e., agree on who will do what), and they can *independently* compute the best path to the task. In our planning algorithm, we use this observation to decompose the decision-making process into a hierarchical structure with two levels:

- At the top level, a task allocation algorithm is run for the whole team to assign the best task to each responder given the current state of the world.

Algorithm 1: Team Coordination

Input: the MMDP model and the current state s.
Output: the best joint action $\vec{a}$.

  // The task planning
1 $\{t_i\} \leftarrow$ compute the best task for each responder $i \in I$
2 **foreach** $i \in I$ **do**
  | // The path planning
3 | $a_i \leftarrow$ compute the best path to task t_i
4 **return** $\vec{a}$

- At the bottom level, given a task, a path planning algorithm is run for each responder to find the best path to the task from his or her current location.

Furthermore, not all states are relevant to the problem (e.g., if a responder gets injured, he or she is incapable of doing any task in the future and therefore his or her states are irrelevant to other responders) and we only need to consider the reachable states given the current global state S of the problem. Hence, given the current state, we compute the policy online only for reachable states. This saves a lot of computation because the size of the reachable states is usually much smaller than the overall state space. For example, given the current location of a responder, the one-step reachable locations are the 8 neighbouring locations plus the current locations, which are 9 locations out of the 50×55 grid. Jointly, the reduction will be huge, which is 9^8 vs. $(50 \times 55)^8$ for 8 responders. Another advantage of online planning is that it allows us to tweak the model as more information is obtained or unexpected events happen. For example, if the wind speed increases or the direction of wind increases, the uncertainty about the radioactive cloud may increase. If a responder becomes tired, the outcome of his or her actions may be liable to more uncertainty.

The main process of our online hierarchical planning algorithm is outlined in Algorithm 1. The following sections will describe the procedures of each level in more detail.

3.1 Task Allocation

As described in Section **??**, each responder p_i has a specific role $r_i \in Roles$ to determine which task he or she can perform and a task t can only be completed by a team of responders with the required roles $Roles(t)$. If, at some point in the execution of a plan, a responder p_i is incapable of performing a task (e.g., because she is tired or suffered a high radiation does), he or she will be removed from the set of responders under consideration that is $I = I \setminus p_i$ if pi. This information can be obtained from the state $s \in S$. When a task is completed by a chosen coalition, the task is simply removed from the set, that is $T = T \setminus t_k$ if t_k has been completed.

Now, in order to capture the efficiency of groupings of responders at performing tasks, we define the notion of a coalition C as a subset of responders, that is, $C \subseteq I$.[1] Thus, we can identify all possible coalitions $\{C_{jk}\}$ for each task t_j where $\{r_i | p_i \in C_{jk}\} = Roles(t_j)$. Crucially, we define the value of a coalition $v(C_{jk})$ that reflects the level of performance of a coalition C_k in performing task t_k. Then, the goal of task allocation algorithm is to assign a task to

---
[1]Here coalitions are not considered in the game-theoretic sense as all agents and coalitions aim to maximise the global objective.

---

each coalition that maximises the overall team performance given the current state s, i.e., $\sum_{j=1}^{m} v(C_j)$ where C_j is a coalition for task t_j and $\{C_1, \cdots, C_m\}$ is a *partition* of I ($\forall j \neq j', C_j \bigcap C_{j'} = \emptyset$ and $\bigcup_{j=1}^{m} C_j = I$). In what follows, we first detail the procedure to compute the value of all coalitions that are valid in a given state and then proceed to detail the main algorithm to allocate tasks. Note that these algorithms take into account the uncertainties captured by the MMDP.

3.1.1 Coalition Value Calculation

The computation of values $v(C_{jk})$ for each coalition C_{jk} is challenging because not all tasks can be completed by one allocation (there are usually more tasks than responders) and the policy after completing task t_j must be computed as well and this is time-consuming. Here, we propose to estimate the value through several simulations. This is much cheaper computationally because we do not need to compute the complete policy in order to come up with a good estimate of the value of the coalition. According to the central limit theorem, as long as the number of simulations is sufficient large, the estimated value will converge to the true coalition value. The main process is outlined in Algorithm 2.

Algorithm 2: Coalition Value Calculation

Input: the current state s, a set of unfinished tasks T, and a set of free responders I.
Output: a task assignment for all responders.

1 $\{C_{jk}\} \leftarrow$ compute all possible coalitions of I for T
2 **foreach** $C_{jk} \in \{C_{jk}\}$ **do**
  | // The N trial simulations
3 | **for** $i = 1$ **to** N **do**
4 | | $(r, s') \leftarrow$ simulate the process with the starting state s until task k is completed by the responders in C_{jk}
5 | | **if** s' *is a terminal state* **then**
6 | | | $v_i(C_{jk}) \leftarrow r$
7 | | **else**
8 | | | $V(s') \leftarrow$ estimate the value of s' with MCTS
9 | | | $v_i(C_{jk}) \leftarrow r + \gamma V(s')$
10 | $v(C_{jk}) \leftarrow \frac{1}{N} \sum_{i=1}^{N} v_i(C_{jk})$
11 **return** the task assignment computed by Equation 3

In each simulation of Algorithm 2, we first assign the responders in C_{jk} to task t_j and run the simulator starting from the current state s (Line 4). After task t_j is completed, the simulator returns the sum of the rewards r and the new state s' (Line 4). If all the responders in C_{jk} are incapable of doing other tasks (e.g., having received too high a radioactive dose), the simulation is terminated (Line 5). Otherwise, we estimate the expected value of s' using Monte-Carlo Tree Search (MCTS) (Line 8), which provides good tradeoff between exploitation and exploration of the policy space and has been shown to be efficient for large MDPs [**?**]. After N simulations, the averaged value is returned as an approximation of the coalition value (Line 10).

The basic idea of MCTS is to maintain a search tree where each node is associated with a state s and each branch is a task assignment for all responders. To implement MCTS, the main step is to compute an assignment for the free responders (a responder is free when he or she is capable of doing tasks but not assigned to any task) at each node of the search tree. This can be computed by Equation 3 using

the coalition values estimated by the UCT heuristic [?]:

$$v(C_{jk}) = \overline{v(C_{jk})} + c\sqrt{\frac{2N(s)}{N(s, C_{jk})}} \qquad (2)$$

where $\overline{v(C_{jk})}$ is the averaged value of coalition C_{jk} at state s so far, c is a tradeoff constant, $N(s)$ is the visiting frequency of state s, and $N(s, C_{jk})$ is the frequency that coalition C_{jk} has been selected at state s. Intuitively, if a coalition C_{jk} has bigger averaged value $\overline{v(C_{jk})}$ or is rarely selected ($N(s, C_{jk})$ is smaller), it has higher chance to be selected in the next visit of the tree node.

As we assume that the role of a responder and the role requirements of each task is static, we can compute all possible coalition values offline and therefore, in the online phase, we only need to filter out the coalitions for completed tasks and those containing incapacitated responders to compute the coalition set $\{C_{jk}\}$.

3.1.2 Coalitional Task Allocation

Given the coalition values computed above, we then solve the following optimisation problem to find the best solution:

$$\begin{aligned} \max_{x_{jk}} \quad & \sum_{j,k} x_{jk} \cdot v(C_{jk}) \\ \text{s.t.} \quad & x_{jk} \in \{0, 1\} \\ & \forall j, \sum_k x_{jk} \le 1 \qquad \text{(i)} \\ & \forall i, \sum_{j,k} \delta_i(C_{jk}) \le 1 \quad \text{(ii)} \end{aligned} \qquad (3)$$

where x_{jk} is the boolean variable to indicate whether coalition C_{jk} is selected for task t_j or not, $v(C_{jk})$ is the characteristic function for coalition C_{jk}, and $\delta_i(C_{jk}) = 1$ if responder $p_i \in C_{jk}$ and 0 otherwise. In the optimisation, Constraint (i) ensures that a task t_j is allocated at most to only one coalition (a task does not need more than one group of responders). Constraint (ii) ensures that a responder p_i is assigned to only one task (a responder cannot do more than one task at the same time). This is a standard MILP that can be efficiently solved using standard solvers such as IBM ILOG's CPLEX.

3.1.3 Adapting to Responders

One main advantage of our approach is that it can easily incorporate the preferences of the responders. For example, if a responder rejects a task allocated to it by the planning agent, we simply filter out the coalitions for the task that contain the responder. By so doing, the responder will not be assigned to the task. Moreover, if a responder prefers to do the tasks with another responder, we can increase the weights of the coalitions that contain them in Equation 3 (By default, all coalitions have identical weights of 1.0). Thus, our approach is adaptive to various preferences of human responders. In particular, we show how the adaptive capability of our algorithm is used in AtomicOrchid in a real-world deployment (in Section ??) Next we show how the path of each responder is computed taking into account real-world uncertainties.

3.2 Path planning

In the path planning phase, we compute the best path for a responder to her assigned task. This phase is stochastic as there are uncertainties in the radioactive cloud and the responders' actions. We model this problem as a single-agent MDP that can be defined as a tuple, $\mathcal{M}_i = \langle S_i, A_i, P_i, R_i \rangle$,

where: $S_i = S_r \times S_{p_i}$ is the state space. In this level, responder p_i only need to consider the states of the radioactive cloud S_r and his or her own states S_{p_i} in the MMDP. A_i is the set of p_i's actions. In this level, responder p_i only need to consider her moving actions. $P_i = P_r \times P_{p_i}$ is the transition function. In this level, responder p_i only need to consider the spreading of the radioactive cloud P_r and the changes of his or her locations and health levels when moving in the filed P_{p_i}, which are defined earlier in the MMDP. R_i is the reward function. At this level, responder p_i only needs to consider the cost of moving to a task and the penalty of receiving high radiation doses.

This is a typical MDP that can be solved by many state-of-the-art MDP solvers [?]. We choose the Real-Time Dynamic Programming (RTDP) [?] approach because it particularly fits our problem, that is, a goal-directed MDP with large number of states. Instead of exploring the whole state space, RTDP only visits the states that are reachable from the initial state s^0 (the start location of the responder). The main process is outline in Algorithm 3. We first compute the Q value for each action (Line 3-4) and select the best action (Line 5). Then, we update the value function (Line 6) and simulate the next state (Line 7). This process repeats until the goal state is reached (Line 9). If the goal is not reached, we assume there does not exist a path between the start location of the responder and the location of the task (either there are obstacles on the path or the responder will be killed by the radioactivity on the road).

Algorithm 3: Path Planning

Input: the starting state s^0 and the goal state s^g.
Output: a path from the starting location to the goal.

1   $s \leftarrow s^0$
2   **repeat**
3   **foreach** $a \in A_i$ **do**
4        $Q(s, a) \leftarrow R_i(s, a) + \sum_{s' \in S_i} P_i(s'|s, a)V(s')$
5      $a \leftarrow \arg\max_{a' \in A_i} Q(s, a')$
6      $V(s) \leftarrow Q(s, a)$
7      $s' \sim P_i(s'|s, a)$
8      $s \leftarrow s'$
9   **until** $s = s^g$
10 **return** Q

There are several techniques we use to speed up the convergency of RTDP. In our problem, the terrain of the field is static. Thus, we can initialize the value function $V(s)$ using the cost of the shortest path between the current location to the task location on the map, which can be computed offline without considering the radioactive cloud. This helps RTDP quickly navigate among the obstacles (e.g., buildings, water pools, blocked roads) without getting trapped in dead-ends during the search. Another speed up is also possible if, when traversing the reachable states (i.e., $s' \in S_i$ in Line 4), we only consider the responder's current location and the neighbouring points, since $P_i(s'|s, a) = 0$ for other points. This will further speed up the algorithm where the main bottleneck is the huge state space.

(**Feng: just provide a graph depicting the performance of the algorithm compared to greedy and say more extensive evaluations are beyond the scope of this paper as the focus is on the use of the algo-**

rithm in a real-world deployment to test how humans take up advice computed in sophisticated ways by an agent-based planner).

4. THE ATOMIC ORCHID PLATFORM

In this section we describe the platform within which we embed the planning agent in order study the interactions between human responders and the agent and derive design guidelines for the implementation of such planning agents in real-world scenarios. (**Joel: please add justification for a mixed-reality game approach to testing this scenario v/s other approaches**).

In more detail, AtomicOrchid is a location-based mobile game based on the fictitious scenario described in Section **??**. Field responders are assigned a specific role (e.g. 'medic', 'transporter', 'soldier', 'ambulance') In their mission to rescue all the targets from the radioactive zone, the field responders are supported by (at least one) person in a centrally located HQ room, and the planning agent that sends the next task (as computed in the previous section) to the team of field responders. In what follows, we first present the player interfaces used, the interactions with the planning agent, and the modelling of the radiation cloud in the game.

4.1 Player interfaces

Field responders are equipped with a 'mobile responder tool' providing sensing and awareness capabilities in three tabs (geiger counter, map, messaging and tasks; see figure XX). One tab shows a reading of radioactivity, player health level (based on exposure), and a GPS-enabled map of the game area to locate fellow responders, the targets to be rescued and the drop off zones for the targets. Another tab provides a broadcast messaging interface to communicate with fellow responders (field responders and HQ). Another tab shows the team and task allocation dynamically provided by the agent. Notifications are used to alert both to new messages and task allocations.

The HQ is manned by at least one player who has at her disposal an 'HQ dashboard' that provides an overview of the game area, including real-time information of the players' locations (see figure XX). The dashboard provides a broadcast messaging widget, and a player status widget so that the responders' exposure and health levels can be monitored. HQ can further monitor the current team and task allocations by the agent. Importantly, only the HQ has a view of the radioactive cloud, depicted as a heatmap. 'Hotter' zones correspond with higher levels of radioactivity.

4.2 System architecture

AtomicOrchid is based on the open-sourced geo-fencing game MapAttack[2] that has been iteratively developed for a responsive, (relatively) scalable experience. The location-based game is realized by client-server architecture, relying on real-time data streaming between client and server.

The client-server architecture is depicted in figure XX. Client-side requests for for less dynamic content use HTTP. Frequent events, such as location updates and radiation exposure, are streamed to clients to avoid the overhead of HTTP. In this way, field responders are kept informed in near real-time.

_____
[2] http://mapattack.org

The platform is built using the geoloqi platform, Sinatra for Ruby, and state-of-the-art web technologies such as socket.io, node.js and the Google Maps API. Open source mobile client apps that are part native, part browser based exist for iPhone and Android; we adapted an Android app to build the mobile responder app.

4.3 Planning agent

[Wenchao. Describe how the agent works (not implementation detail, add that in subsection below), i.e., when it is polled, what information is being exchanged, and how the team/task allocation is being constructed from that and sent.] The planning agent is a standalone software agent designed to solve coordination problems in the AtomicOrhid game scenario. It takes game status as input and calculate solutions by running xxx Algorithms.

4.4 Integration with AtomicOrchid

The agent is implemented by JAVA [need Feng's confirmation] and deployed on a server separated from AtomicOrchid platform. The agent and AtomicOrhid communicate through a simple HTTP interface. The AtomicOrchid server send HTTP requests to pull plans from planning Agent whenever replans are triggered in game.

There are two triggers of replanning in the game.

- *Successful drop off*. explain more.....

- *Explicit reject*. Field players can explicitly reject a plan by pressing reject button in mobile responder app.

The interactions between agent and players will be detailed in next section.

4.5 Interacting with planning agent

4.6 Radiation Cloud Modelling

The radiation cloud is assumed to be monitored using a number of sensors on the ground (within the disaster space) that collect readings of the radiation cloud intensity and wind velocity every minute of the game. These sensors can be at fixed locations or held by mobile agents. The radiation cloud diffusion process is modelled by a nonlinear Markov field stochastic differential equation,

$$\frac{\partial \rho(\mathbf{r}, t)}{\partial t} = \kappa \nabla^2 \rho(\mathbf{r}, t) + \nabla \rho(\mathbf{r}, t).\mathbf{v}(\mathbf{r}, t)$$

where ρ is the radiation cloud intensity at location $\mathbf{r}$ at time t and κ is a fixed duffusion coefficient. The diffusion equation is solved on a regular grid defined across the environment with grid coordinates G. Furthermore, the grid is solved at discrete time instances t. The cloud is driven by wind forces which vary both spatially and temporally. These forces induce anisotropy into the cloud diffusion process which is proportional to the wind velocity, $\mathbf{v}(\mathbf{r},t)$. The wind velocity is drawn from two independent Gaussian processes (GP), one GP for each Cartesian coordinate axis, $v_i(\mathbf{r}, t)$, of $\mathbf{v}(\mathbf{r}, t)$. The GP captures both the spatial distribution of the wind velocity and also the dynamic process resulting from shifting wind patterns such as short term gusts and longer term variations. In our simulation, each spatial wind velocity component is modelled by a squared-exponential GP covariance function, K, with fixed input and output scales (although any covariance function can be substituted). Furthermore,

as wind conditions may change over time we introduce a temporal correlation coefficient ρ to the covariance function. Thus, for a single component, v_i, of $\mathbf{v}$, defined over grid G at times t and t', the wind process covariance function is, $\text{Cov}(v_i(G, t), v_i(G, t')) = \rho(t, t')K(G, G)$. We note that, when $\rho = 1$ the wind velocities are time invariant (although spatially variant). Values of $\rho < 1$ model wind conditions that change over time.

(**Steve: in the platform we take the 'real' values from the diffusion process i believe. Does the above capture this? We will say that we will add the features you mention below to a future version of the platform where we aim to do both situational awareness and rescue. Add a sentence above to conclude where we took the values from and the process takes into account the location of radiation source. Also, your notation clashes with the notations in the scenario and Feng's algorithm - please try to align.**)

5. REAL-WORLD EVALUATION

[Note: currently not sure whether to include the non-agent runs. Problematic because: a) unequal number of responders, b) HQ staffed by students in non-agent condition; researchers in agent-condition, c) not enough cases for a quantitative comparison anyways?.] (**Joel: the population type does not matter so much for AAMAS as tar as I've seen. So go ahead with the analysis of the no-agent condition.**) We ran four sessions of AtomicOrchid with participants recruited from the local university to evaluate mixed-initiative coordination in a disaster response scenario. The following sections describe the participants, procedure, session configuration and methods used to collect and analyse quantitative and qualitative data.

5.1 Participants

A total of 29 participants (XX of them were female) were recruited through posters and emails, and reimbursed with 15 pounds for 1.5-2 hours of study. The majority were students of the local university. [Say something about their map reading skills?]

5.2 Procedure

The procedure consisted of 30 minutes of game play, and about 1 hour of pre-game briefing, consent forms and a short training session, and post-game group discussion and questionnaire.

At the end of the briefing in which mission objectives and rules were outlined, responder roles were randomly assigned to all participants (fire-fighter, medic, transporter, soldier). HQ in the agent condition was staffed by a different member of the research team in each session in order to mimic an experienced HQ whilst avoiding the same person running HQ every time.

Field responders were provided with a smartphone; HQ coordinators with a laptop. The team was given 5 minutes to discuss a common game strategy. (**Joel: where did the agent run ?**)

Field responders were then accompanied to the starting point within the designated game area, about 1 minute walk from headquarters. Once field responders were ready to start, HQ sent a 'game start' message. After 30 minutes of game play the field responders returned to the HQ for the post-game session, which consisted of a questionnaire aimed at collecting participants' feedback on (1) first impressions of the game; (2) usability of the system, and; (3) coordination issues in the game. A group interview was then conducted, before participants were debriefed and dismissed.

5.3 Game sessions

We ran two sessions without the planner agent, and two sessions with the planner agent to be able to compare team performance in the two conditions. We also ran a pilot study for each condition. The pilot study showed that this was a challenging, yet not too overwhelming number of targets to collect in a 30 min game session. There were four targets for each of the four target types. The target locations, pattern of cloud movement and expansion were kept constant for all game sessions.

The role allocation of the 8 field responders per session is depicted in table XX. One of the non-agent sessions only had 5 field responders due to drop outs.

The terrain of the game area includes grassland, a lake, buildings, roads, and footpaths and lawns (see figure XX). There are two drop off zones and 16 targets.

5.4 Methods

We took a mixed methods approach to data collection and analysis. In addition to quantitative questionnaires, a semi-structured group interview was conducted that aimed at eliciting important decision points, strategies and the overall decision-making process. Furthermore, researchers with camcorders recorded the game play. One researcher recorded action in the HQ, and four other researchers each shadowed a field responder team with a camcorder.

We developed a log file replay tool to help with data analysis of time stamped system logs that contain a complete record of the game play, including responders' GPS location, their health status and radioactive exposure, messages, cloud location, locations of target objects and task status.

Video recordings of field action were catalogued to identify sequences (episodes) of interest (cf. Heath et al., 2010). Key decision points in teaming and task allocation served to index the episodes. Interesting distinct units of interaction were transcribed and triangulated with log files of relevant game activity for deeper analysis. Due to space constraints we can only present one fragment in this paper to illustrate how human-agent collaboration typically unfolded (TODO).

How are remote messages used as a coordination resource? We use speech-act theory (Searle, 1975) to classify messages sent between and among responders and HQ. We focus on the most relevant types of acts in this paper (which are also the most frequently used in AtomicOrchid):

- Assertives: *speech acts that commit a speaker to the truth of the expressed proposition*; these were a common category as they include messages that contain situational information.

- Directives: *speech acts that are meant to cause the hearer to take a particular action*, e.g. requests, commands and advice, including task and team allocation messages.

5.5 Results

Structure.

- *Overall performance.* Draw on metrics below: tasks completed, number and categorisation of messages (only directives and assertives).

  task finished in 1.43 in agent-enabled session [session a 8] task finished in 1.43 in agent-enabled session [session b 12, session c 11] average number of team reformation are 12.5 in agent-enabled session [session b 4 , session c 9]

  The avg health in agent condition is 80.75 (session b 89.75, session c 71.75 The avg health in non-agent condition is 39.695

- *Agent performance.* Metrics from below: Number of instructions sent, robustness etc. Replan happened 32 times across 2 agent-enabled sessions [14,18]. Triggered by rejection 8 times, triggered by drop off 24 times [1 unsuccessful drop off, 23 successful drop off]

  [2 diagrams are ready agent teaming instruction, HQ directives classification, observed 2 violations of agent planning].

- *Task allocation*: How task allocation unfolded in the agent vs. non-agent condition. (Message handling (from JSCWS paper) vs. task handling diagram...) This is where we'd show a fragment to illustrate? -> Shows overall performance increase in performance [begin looking for fragments now]

- *Rejecting tasks*: When and why did it happen? (-> pick this up in the discussion re. Gopal's/Feng's point on adjustable planning?). Players are more likely to reject plan if their proposed teammates are far away from them. For accepted instructions, the average distance between related players are 12 meters. For rejected instructions, the average distance between related players are 86 meters.

- *The role of HQ*: monitoring, supporting and dealing with contingencies. Some example messages. Draw on HQ metrics. (-> Shows division of labour and the benefits of human-agent collaboration). HQ have to click "view" button to reveal agent plan. We observed 57 clicks in the 2 agent-enabled sessions.

Joel and Wenchao

1. Explain setup of experiment - area of interest + setup of tasks

2. Explain evaluation = quantitative and qualitative.

Metrics.

- Comparisons between with/without agent versions for the below:

- Performance of FR: number of tasks completed, time on task?, number of messages sent, number of teams formed and disbanded, time on team, acknowledgements of tasks

- Messages: classification

- Health

- Distance travelled

- HQ: number of agent monitoring actions (clicks), number of 'supporting'/related messages (e.g., enforcement, contradictions/overriding)

- Agent performance: number of instructions, number of replanning steps, replanning robustness (diversion of task allocation compared to previous step)

- Following instructions ('obedience'): number of instructions followed vs. not followed (incl. number of HQ interventions/overriding agent allocation), instruction handling diagram

-

5.6   Conclusions

6.   REFERENCES